

6.824 2020 Lecture 2: Infrastructure: RPC and threads

Today:

Threads and RPC in Go, with an eye towards the labs

Why Go?

- good support for threads
- convenient RPC
- type safe
- garbage-collected (no use after freeing problems)
- threads + GC is particularly attractive!
- relatively simple

After the tutorial, use https://golang.org/doc/effective_go.html

Threads

- a useful structuring tool, but can be tricky
- Go calls them goroutines; everyone else calls them threads

Thread = "thread of execution"

- threads allow one program to do many things at once
- each thread executes serially, just like an ordinary non-threaded program
- the threads share memory
- each thread includes some per-thread state:
 - program counter, registers, stack

Why threads?

- They express concurrency, which you need in distributed systems
- I/O concurrency
 - Client sends requests to many servers in parallel and waits for replies.
 - Server processes multiple client requests; each request may block.
 - While waiting for the disk to read data for client X,
 - process a request from client Y.
- Multicore performance
 - Execute code in parallel on several cores.
- Convenience
 - In background, once per second, check whether each worker is still alive.

Is there an alternative to threads?

- Yes: write code that explicitly interleaves activities, in a single thread.
 - Usually called "event-driven."
- Keep a table of state about each activity, e.g. each client request.
- One "event" loop that:
 - checks for new input for each activity (e.g. arrival of reply from server),
 - does the next step for each activity,
 - updates state.
- Event-driven gets you I/O concurrency,
 - and eliminates thread costs (which can be substantial),
 - but doesn't get multi-core speedup,
 - and is painful to program.

Threading challenges:

- shared data
 - e.g. what if two threads do $n = n + 1$ at the same time?
 - or one thread reads while another increments?
 - this is a "race" -- and is usually a bug
 - > use locks (Go's `sync.Mutex`)
 - > or avoid sharing mutable data
- coordination between threads
 - e.g. one thread is producing data, another thread is consuming it
 - how can the consumer wait (and release the CPU)?
 - how can the producer wake up the consumer?
 - > use Go channels or `sync.Cond` or `WaitGroup`
- deadlock
 - cycles via locks and/or communication (e.g. RPC or Go channels)

Let's look at the tutorial's web crawler as a threading example.

What is a web crawler?

- goal is to fetch all web pages, e.g. to feed to an indexer
- web pages and links form a graph
- multiple links to some pages
- graph has cycles

Crawler challenges

- Exploit I/O concurrency
 - Network latency is more limiting than network capacity
 - Fetch many URLs at the same time
 - To increase URLs fetched per second
 - => Need threads for concurrency
- Fetch each URL only *once*
 - avoid wasting network bandwidth
 - be nice to remote servers
 - => Need to remember which URLs visited
- Know when finished

We'll look at two styles of solution [crawler.go on schedule page]

Serial crawler:

- performs depth-first exploration via recursive Serial calls
- the "fetched" map avoids repeats, breaks cycles
 - a single map, passed by reference, caller sees callee's updates
- but: fetches only one page at a time
 - can we just put a "go" in front of the Serial() call?
 - let's try it... what happened?

ConcurrentMutex crawler:

- Creates a thread for each page fetch
 - Many concurrent fetches, higher fetch rate
- the "go func" creates a goroutine and starts it running
 - func... is an "anonymous function"
- The threads share the "fetched" map
 - So only one thread will fetch any given page
- Why the Mutex (Lock() and Unlock())?
 - One reason:
 - Two different web pages contain links to the same URL
 - Two threads simultaneously fetch those two pages
 - T1 reads fetched[url], T2 reads fetched[url]
 - Both see that url hasn't been fetched (already == false)
 - Both fetch, which is wrong
 - The lock causes the check and update to be atomic
 - So only one thread sees already==false
 - Another reason:
 - Internally, map is a complex data structure (tree? expandable hash?)
 - Concurrent update/update may wreck internal invariants
 - Concurrent update/read may crash the read
- What if I comment out Lock() / Unlock()?
 - go run crawler.go
 - Why does it work?
 - go run -race crawler.go
 - Detects races even when output is correct!
- How does the ConcurrentMutex crawler decide it is done?
 - sync.WaitGroup
 - Wait() waits for all Add()s to be balanced by Done()s
 - i.e. waits for all child threads to finish
 - [diagram: tree of goroutines, overlaid on cyclic URL graph]
 - there's a WaitGroup per node in the tree
- How many concurrent threads might this crawler create?

ConcurrentChannel crawler

```

a Go channel:
  a channel is an object
    ch := make(chan int)
  a channel lets one thread send an object to another thread
  ch <- x
    the sender waits until some goroutine receives
  y := <- ch
    for y := range ch
      a receiver waits until some goroutine sends
  channels both communicate and synchronize
  several threads can send and receive on a channel
  channels are cheap
  remember: sender blocks until the receiver receives!
    "synchronous"
  watch out for deadlock
ConcurrentChannel master()
  master() creates a worker goroutine to fetch each page
  worker() sends slice of page's URLs on a channel
    multiple workers send on the single channel
  master() reads URL slices from the channel
At what line does the master wait?
  Does the master use CPU time while it waits?
No need to lock the fetched map, because it isn't shared!
How does the master know it is done?
  Keeps count of workers in n.
  Each worker sends exactly one item on channel.

```

Why is it not a race that multiple threads use the same channel?

Is there a race when worker thread writes into a slice of URLs,
and master thread reads that slice, without locking?

- * worker only writes slice **before** sending
- * master only reads slice **after** receiving

So they can't use the slice at the same time.

When to use sharing and locks, versus channels?

Most problems can be solved in either style
What makes the most sense depends on how the programmer thinks

- state -- sharing and locks
- communication -- channels

For the 6.824 labs, I recommend sharing+locks for state,
and sync.Cond or channels or time.Sleep() for waiting/notification.

Remote Procedure Call (RPC)

a key piece of distributed system machinery; all the labs use RPC
goal: easy-to-program client/server communication
hide details of network protocols
convert data (strings, arrays, maps, &c) to "wire format"

RPC message diagram:

```

Client          Server
request--->
<---response

```

Software structure

```

client app      handler fns
stub fns        dispatcher
RPC lib         RPC lib
net  ----- net

```

Go example: kv.go on schedule page

A toy key/value storage server -- Put(key,value), Get(key)->value
Uses Go's RPC library
Common:

Declare Args and Reply struct for each server handler.

Client:

- connect()'s Dial() creates a TCP connection to the server
- get() and put() are client "stubs"
- Call() asks the RPC library to perform the call
 - you specify server function name, arguments, place to put reply
 - library marshalls args, sends request, waits, unmarshalls reply
 - return value from Call() indicates whether it got a reply
 - usually you'll also have a reply.Err indicating service-level failure

Server:

- Go requires server to declare an object with methods as RPC handlers
- Server then registers that object with the RPC library
- Server accepts TCP connections, gives them to RPC library
- The RPC library
 - reads each request
 - creates a new goroutine for this request
 - unmarshalls request
 - looks up the named object (in table created by Register())
 - calls the object's named method (dispatch)
 - marshalls reply
 - writes reply on TCP connection
- The server's Get() and Put() handlers
 - Must lock, since RPC library creates a new goroutine for each request
 - read args; modify reply

A few details:

- Binding: how does client know what server computer to talk to?
 - For Go's RPC, server name/port is an argument to Dial
 - Big systems have some kind of name or configuration server
- Marshalling: format data into packets
 - Go's RPC library can pass strings, arrays, objects, maps, &c
 - Go passes pointers by copying the pointed-to data
 - Cannot pass channels or functions

RPC problem: what to do about failures?

- e.g. lost packet, broken network, slow server, crashed server

What does a failure look like to the client RPC library?

- Client never sees a response from the server
- Client does **not** know if the server saw the request!
 - [diagram of losses at various points]
 - Maybe server never saw the request
 - Maybe server executed, crashed just before sending reply
 - Maybe server executed, but network died just before delivering reply

Simplest failure-handling scheme: "best effort"

- Call() waits for response for a while
- If none arrives, re-send the request
- Do this a few times
- Then give up and return an error

Q: is "best effort" easy for applications to cope with?**A particularly bad situation:**

- client executes
 - Put("k", 10);
 - Put("k", 20);
- both succeed
- what will Get("k") yield?
 - [diagram, timeout, re-send, original arrives late]

Q: is best effort ever OK?

- read-only operations
- operations that do nothing if repeated
 - e.g. DB checks if record has already been inserted

Better RPC behavior: "at most once"

- idea: server RPC code detects duplicate requests
- returns previous reply instead of re-running handler

Q: how to detect a duplicate request?

- client includes unique ID (XID) with each request

- uses same XID for re-send

server:

```
if seen[xid]:
    r = old[xid]
else
    r = handler()
    old[xid] = r
    seen[xid] = true
```

some at-most-once complexities

- this will come up in lab 3

- what if two clients use the same XID?

- big random number?

- combine unique client ID (ip address?) with sequence #?

- server must eventually discard info about old RPCs

- when is discard safe?

- idea:

- each client has a unique ID (perhaps a big random number)

- per-client RPC sequence numbers

- client includes "seen all replies $\leq X$ " with every RPC

- much like TCP sequence #s and acks

- or only allow client one outstanding RPC at a time

- arrival of seq+1 allows server to discard all \leq seq

- how to handle dup req while original is still executing?

- server doesn't know reply yet

- idea: "pending" flag per executing RPC; wait or ignore

What if an at-most-once server crashes and re-starts?

- if at-most-once duplicate info in memory, server will forget

- and accept duplicate requests after re-start

- maybe it should write the duplicate info to disk

- maybe replica server should also replicate duplicate info

Go RPC is a simple form of "at-most-once"

- open TCP connection

- write request to TCP connection

- Go RPC never re-sends a request

- So server won't see duplicate requests

- Go RPC code returns an error if it doesn't get a reply

- perhaps after a timeout (from TCP)

- perhaps server didn't see request

- perhaps server processed request but server/net failed before reply came back

What about "exactly once"?

- unbounded retries plus duplicate detection plus fault-tolerant service

- Lab 3