

6.824 2020 Lecture 4: Primary/Backup Replication

Today

Primary/Backup Replication for Fault Tolerance
Case study of VMware FT, an extreme version of the idea

The topic is (still) fault tolerance
to provide availability
despite server and network failures
using replication

What kinds of failures can replication deal with?

"fail-stop" failure of a single replica
fan stops working, CPU overheats and shuts itself down
someone trips over replica's power cord or network cable
software notices it is out of disk space and stops
Maybe not defects in h/w or bugs in s/w or human configuration errors
Often not fail-stop
May be correlated (i.e. cause all replicas to crash at the same time)
But, sometimes can be detected (e.g. checksums)
How about earthquake or city-wide power failure?
Only if replicas are physically separated

Is replication worth the Nx expense?

Two main replication approaches:

State transfer
Primary replica executes the service
Primary sends [new] state to backups
Replicated state machine
Clients send operations to primary,
primary sequences and sends to backups
All replicas execute all operations
If same start state,
same operations,
same order,
deterministic,
then same end state.

State transfer is simpler
But state may be large, slow to transfer over network

Replicated state machine often generates less network traffic
Operations are often small compared to state
But complex to get right
VM-FT uses replicated state machine, as do Labs 2/3/4

Big Questions:

What state to replicate?
Does primary have to wait for backup?
When to cut over to backup?
Are anomalies visible at cut-over?
How to bring a replacement backup up to speed?

At what level do we want replicas to be identical?

Application state, e.g. a database's tables?
GFS works this way
Can be efficient; primary only sends high-level operations to backup
Application code (server) must understand fault tolerance, to e.g. forward op stream
Machine level, e.g. registers and RAM content?
might allow us to replicate any existing server w/o modification!
requires forwarding of machine events (interrupts, DMA, &c)
requires "machine" modifications to send/recv event stream...

Today's paper (VMware FT) replicates machine-level state
 Transparent: can run any existing O/S and server software!
 Appears like a single server to clients

Overview

[diagram: app, O/S, VM-FT underneath, disk server, network, clients]

words:

hypervisor == monitor == VMM (virtual machine monitor)
 O/S+app is the "guest" running inside a virtual machine
 two machines, primary and backup
 primary sends all external events (client packets &c) to backup over network
 "logging channel", carrying log entries
 ordinarily, backup's output is suppressed by FT
 if either stops being able to talk to the other over the network
 "goes live" and provides sole service
 if primary goes live, it stops sending log entries to the backup

VMM emulates a local disk interface

but actual storage is on a network server

treated much like a client:

usually only primary communicates with disk server (backup's FT discards)
 if backup goes live, it talks to disk server
 external disk makes creating a new backup faster (don't have to copy primary's disk)

When does the primary have to send information to the backup?

Any time something happens that might cause their executions to diverge.
 Anything that's not a deterministic consequence of executing instructions.

What sources of divergence must FT handle?

Most instructions execute identically on primary and backup.

As long as memory+registers are identical,
 which we're assuming by induction.

Inputs from external world -- just network packets.

These appear as DMA'd data plus an interrupt.

Timing of interrupts.

Instructions that aren't functions of state, such as reading current time.

Not multi-core races, since uniprocessor only.

Why would divergence be a disaster?

b/c state on backup would differ from state on primary,
 and if primary then failed, clients would see inconsistency.

Example: GFS lease expiration

Imagine we're replicating the GFS master

Chunkserver must send "please renew" msg before 60-second lease expires

Clock interrupt drives master's notion of time

Suppose chunkserver sends "please renew" just around 60 seconds

On primary, clock interrupt happens just after request arrives.

Primary copy of master renews the lease, to the same chunkserver.

On backup, clock interrupt happens just before request.

Backup copy of master expires the lease.

If primary fails, backup takes over, it will think there
 is no lease, and grant it to a different chunkserver.

Then two chunkservers will have lease for same chunk.

So: backup must see same events,

in same order,

at same points in instruction stream.

Each log entry: instruction #, type, data.

FT's handling of timer interrupts

Goal: primary and backup should see interrupt at
 the same point in the instruction stream

Primary:

FT fields the timer interrupt

FT reads instruction number from CPU

FT sends "timer interrupt at instruction X" on logging channel
 FT delivers interrupt to primary, and resumes it
 (this relies on CPU support to interrupt after the X'th instruction)

Backup:

ignores its own timer hardware
 FT sees log entry *before* backup gets to instruction X
 FT tells CPU to interrupt (to FT) at instruction X
 FT mimics a timer interrupt to backup

FT's handling of network packet arrival (input)

Primary:

FT tells NIC to copy packet data into FT's private "bounce buffer"
 At some point NIC does DMA, then interrupts
 FT gets the interrupt
 FT pauses the primary
 FT copies the bounce buffer into the primary's memory
 FT simulates a NIC interrupt in primary
 FT sends the packet data and the instruction # to the backup

Backup:

FT gets data and instruction # from log stream
 FT tells CPU to interrupt (to FT) at instruction X
 FT copies the data to backup memory, simulates NIC interrupt in backup

Why the bounce buffer?

We want the data to appear in memory at exactly the same point in execution of the primary and backup.

Otherwise they may diverge.

Note that the backup must lag by one one log entry

Suppose primary gets an interrupt, or input, after instruction X
 If backup has already executed past X, it cannot handle the input correctly
 So backup FT can't start executing at all until it sees the first log entry
 Then it executes just to the instruction # in that log entry
 And waits for the next log entry before resuming backup

Example: non-deterministic instructions

some instructions yield different results even if primary/backup have same state
 e.g. reading the current time or cycle count or processor serial #

Primary:

FT sets up the CPU to interrupt if primary executes such an instruction
 FT executes the instruction and records the result
 sends result and instruction # to backup

Backup:

FT reads log entry, sets up for interrupt at instruction #
 FT then supplies value that the primary got

What about output (sending network packets)?

Primary and backup both execute instructions for output
 Primary's FT actually does the output
 Backup's FT discards the output

Output example: DB server

clients can send "increment" request
 DB increments stored value, replies with new value
 so:

[diagram]
 suppose the server's value starts out at 10
 network delivers client request to FT on primary
 primary's FT sends on logging channel to backup
 FTs deliver request to primary and backup
 primary executes, sets value to 11, sends "11" reply, FT really sends reply
 backup executes, sets value to 11, sends "11" reply, and FT discards
 the client gets one "11" response, as expected

But wait:

suppose primary crashes just after sending the reply
 so client got the "11" reply
 AND the logging channel discards the log entry w/ client request
 primary is dead, so it won't re-send
 backup goes live
 but it has value "10" in its memory!
 now a client sends another increment request
 it will get "11" again, not "12"
 oops

Solution: the Output Rule (Section 2.2)

before primary sends output,
 must wait for backup to acknowledge all previous log entries

Again, with output rule:

[diagram]
 primary:
 receives client "increment" request
 sends client request on logging channel
 about to send "11" reply to client
 first waits for backup to acknowledge previous log entry
 then sends "11" reply to client
 suppose the primary crashes at some point in this sequence
 if before primary receives acknowledgement from backup
 maybe backup didn't see client's request, and didn't increment
 but also primary won't have replied
 if after primary receives acknowledgement from backup
 then client may see "11" reply
 but backup guaranteed to have received log entry w/ client's request
 so backup will increment to 11

The Output Rule is a big deal

Occurs in some form in all replication systems
 A serious constraint on performance
 An area for application-specific cleverness
 Eg. maybe no need for primary to wait before replying to read-only operation
 FT has no application-level knowledge, must be conservative

Q: What if the primary crashes just after getting ACK from backup,
 but before the primary emits the output?
 Does this mean that the output won't ever be generated?

A: Here's what happens when the primary fails and the backup goes live.
 The backup got some log entries from the primary.
 The backup continues executing those log entries WITH OUTPUT DISCARDED.
 After the last log entry, the backup goes live -- stops discarding output
 In our example, the last log entry is arrival of client request
 So after client request arrives, the client will start emitting outputs
 And thus it will emit the reply to the client

Q: But what if the primary crashed *after* emitting the output?
 Will the backup emit the output a *second* time?

A: Yes.
 OK for TCP, since receivers ignore duplicate sequence numbers.
 OK for writes to disk, since backup will write same data to same block #.

Duplicate output at cut-over is pretty common in replication systems
 Clients need to keep enough state to ignore duplicates
 Or be designed so that duplicates are harmless

Q: Does FT cope with network partition -- could it suffer from split brain?
 E.g. if primary and backup both think the other is down.
 Will they both go live?

A: The disk server breaks the tie.

Disk server supports atomic test-and-set.

If primary or backup thinks other is dead, attempts test-and-set.

If only one is alive, it will win test-and-set and go live.

If both try, one will lose, and halt.

The disk server may be a single point of failure

If disk server is down, service is down

They probably have in mind a replicated disk server

Q: Why don't they support multi-core?

Performance (table 1)

FT/Non-FT: impressive!

little slow down

Logging bandwidth

Directly reflects disk read rate + network input rate

18 Mbit/s for my-sql

These numbers seem low to me

Applications can read a disk at at least 400 megabits/second

So their applications aren't very disk-intensive

When might FT be attractive?

Critical but low-intensity services, e.g. name server.

Services whose software is not convenient to modify.

What about replication for high-throughput services?

People use application-level replicated state machines for e.g. databases.

The state is just the DB, not all of memory+disk.

The events are DB commands (put or get), not packets and interrupts.

Result: less fine-grained synchronization, less overhead.

GFS use application-level replication, as do Lab 2 &c

Summary:

Primary-backup replication

VM-FT: clean example

How to cope with partition without single point of failure?

Next lecture

How to get better performance?

Application-level replicated state machines

VMware KB (#1013428) talks about multi-CPU support. VM-FT may have switched from a replicated state machine approach to the state transfer approach, but unclear whether that is true or not.

<http://www.wooditwork.com/2014/08/26/whats-new-vsphere-6-0-fault-tolerance/>

<http://www.mount.ece.umn.edu/~jjyi/MoBS/2007/program/01C-Xu.pdf>