

6.824 2020 Lecture 6: Raft (1)

this lecture

today: Raft elections and log handling (Lab 2A, 2B)

next: Raft persistence, client behavior, snapshots (Lab 2C, Lab 3)

a pattern in the fault-tolerant systems we've seen

- * MR replicates computation but relies on a single master to organize

- * GFS replicates data but relies on the master to pick primaries

- * VMware FT replicates service but relies on test-and-set to pick primary

all rely on a single entity to make critical decisions

nice: decisions by a single entity avoid split brain

how could split brain arise, and why is it damaging?

suppose we're replicating a test-and-set service

the client request sets the state to 1, server replies w/ previous state

only one client should get a reply with "0" !!!

it's a lock, only one requester should get it

[C1, C2, S1, S2]

suppose client C1 can contact replica S1, but not replica S2

should C1 proceed with just replica S1?

if S2 has really crashed, C1 *must* proceed without S2,

otherwise the service doesn't tolerate faults!

if S2 is up but network prevents C1 from contacting S2,

C1 should *not* proceed without S2,

since S2 might be alive and serving client C2

with this setup, we're faced with a nasty choice:

either no ability to tolerate faults, despite replication, or

the possibility of incorrect operation due to split brain

the problem: computers cannot distinguish "server crashed" vs "network broken"

the symptom is the same: no response to a query over the network

the bad situation is often called "network partition":

C1 can talk to S1, C2 can talk to S2,

but C1+S1 see no responses from C2+S2

this difficulty seemed insurmountable for a long time

seemed to require outside agent (a human) to decide when to cut over

or a single perfectly reliable server (FT's test-and-set server)

or a perfectly reliable network (so "no response" == "crashed")

BUT these are all single points of failure -- not desirable

can one do better?

The big insight for coping w/ partition: majority vote

require an odd number of servers, e.g. 3

agreement from a majority is required to do anything -- 2 out of 3

why does majority help avoid split brain?

at most one partition can have a majority

breaks the symmetry we saw with just two servers

note: majority is out of all servers, not just out of live ones

more generally $2f+1$ can tolerate f failed servers

since the remaining $f+1$ is a majority of $2f+1$

if more than f fail (or can't be contacted), no progress

often called "quorum" systems

a key property of majorities is that any two must intersect

e.g. successive majorities for Raft leader election must overlap

and the intersection can convey information about previous decisions

Two partition-tolerant replication schemes were invented around 1990,

Paxos and View-Stamped Replication

in the last 15 years this technology has seen a lot of real-world use

the Raft paper is a good introduction to modern techniques

*** topic: Raft overview

state machine replication with Raft -- Lab 3 as example:

[diagram: clients, 3 replicas, k/v layer + state, raft layer + logs]

Raft is a library included in each replica

time diagram of one client command

[C, L, F1, F2]

client sends Put/Get "command" to k/v layer in leader

leader adds command to log

leader sends AppendEntries RPCs to followers

followers add command to log

leader waits for replies from a bare majority (including itself)

entry is "committed" if a majority put it in their logs

committed means won't be forgotten even if failures

majority -> will be seen by the next leader's vote requests

leader executes command, replies to client

leader "piggybacks" commit info in next AppendEntries

followers execute entry once leader says it's committed

why the logs?

the service keeps the state machine state, e.g. key/value DB

why isn't that enough?

the log orders the commands

to help replicas agree on a single execution order

to help the leader ensure followers have identical logs

the log stores tentative commands until committed

the log stores commands in case leader must re-send to followers

the log stores commands persistently for replay after reboot

are the servers' logs exact replicas of each other?

no: some replicas may lag

no: we'll see that they can temporarily have different entries

the good news:

they'll eventually converge to be identical

the commit mechanism ensures servers only execute stable entries

lab 2 Raft interface

rf.Start(command) (index, term, isleader)

Lab 3 k/v server's Put()/Get() RPC handlers call Start()

Start() only makes sense on the leader

starts Raft agreement on a new log entry

add to leader's log

leader sends out AppendEntries RPCs

Start() returns w/o waiting for RPC replies

k/v layer's Put()/Get() must wait for commit, on applyCh

agreement might fail if server loses leadership before committing

then the command is likely lost, client must re-send

isleader: false if this server isn't the leader, client should try another

term: currentTerm, to help caller detect if leader is later demoted

index: log entry to watch to see if the command was committed

ApplyMsg, with Index and Command

each peer sends an ApplyMsg on applyCh for each committed entry

each peer's local service code executes, updates local replica state

leader sends reply to waiting client RPC

there are two main parts to Raft's design:

electing a new leader

ensuring identical logs despite failures

*** topic: leader election (Lab 2A)

why a leader?

ensures all replicas execute the same commands, in the same order

(some designs, e.g. Paxos, don't have a leader)

Raft numbers the sequence of leaders

new leader -> new term

a term has at most one leader; might have no leader
the numbering helps servers follow latest leader, not superseded leader

when does a Raft peer start a leader election?

when it doesn't hear from current leader for an "election timeout"
increments local currentTerm, tries to collect votes
note: this can lead to un-needed elections; that's slow but safe
note: old leader may still be alive and think it is the leader

how to ensure at most one leader in a term?

(Figure 2 RequestVote RPC and Rules for Servers)
leader must get "yes" votes from a majority of servers
each server can cast only one vote per term
if candidate, votes for itself
if not a candidate, votes for first that asks (within Figure 2 rules)
at most one server can get majority of votes for a given term
-> at most one leader even if network partition
-> election can succeed even if some servers have failed

how does a server learn about newly elected leader?

new leader sees yes votes from majority
others see AppendEntries heart-beats with a higher term number
i.e. from the new leader
the heart-beats suppress any new election

an election may not succeed for two reasons:

- * less than a majority of servers are reachable
- * simultaneous candidates split the vote, none gets majority

what happens if an election doesn't succeed?

another timeout (no heartbeat), a new election (and new term)
higher term takes precedence, candidates for older terms quit

how does Raft avoid split votes?

each server picks a random election timeout
[diagram of times at which servers' timeouts expire]
randomness breaks symmetry among the servers
one will choose lowest random delay
hopefully enough time to elect before next timeout expires
others will see new leader's AppendEntries heartbeats and
not become candidates
randomized delays are a common pattern in network protocols

how to choose the election timeout?

- * at least a few heartbeat intervals (in case network drops a heartbeat)
to avoid needless elections, which waste time
- * random part long enough to let one candidate succeed before next starts
- * short enough to react quickly to failure, avoid long pauses
- * short enough to allow a few re-tries before tester gets upset
tester requires election to complete in 5 seconds or less

what if old leader isn't aware a new leader is elected?

perhaps old leader didn't see election messages
perhaps old leader is in a minority network partition
new leader means a majority of servers have incremented currentTerm
so old leader (w/ old term) can't get majority for AppendEntries
so old leader won't commit or execute any new log entries
thus no split brain
but a minority may accept old server's AppendEntries
so logs may diverge at end of old term