

6.824 2020 Lecture 12: Distributed Transactions

Topics:

distributed transactions = concurrency control + atomic commit

what's the problem?

lots of data records, sharded on multiple servers, lots of clients

[diagram: clients, servers, data sharded by key]

client application actions often involve multiple reads and writes

bank transfer: debit and credit

vote: check if already voted, record vote, increment count

install bi-directional links in a social graph

we'd like to hide interleaving and failure from application writers

this is a traditional database concern

today's material originated with databases

but the ideas are used in many distributed systems

the traditional plan: transactions

programmer marks beginning/end of sequences of code as transactions

example transactions

x and y are bank balances -- records in database tables

x and y are on different servers (maybe at different banks)

x and y start out as \$10

T1 and T2 are transactions

T1: transfer \$1 from x to y

T2: audit, to check that no money is lost

T1:

T2:

begin_xaction begin_xaction

add(x, 1) tmp1 = get(x)

add(y, -1) tmp2 = get(y)

end_xaction print tmp1, tmp2

end_xaction

what is correct behavior for a transaction?

usually called "ACID"

Atomic -- all writes or none, despite failures

Consistent -- obeys application-specific invariants

Isolated -- no interference between xactions -- serializable

Durable -- committed writes are permanent

we're interested in ACID for distributed transactions

with data sharded over multiple servers

What does serializable mean?

you execute some concurrent transactions, which yield results

"results" means both output and changes in the DB

the results are serializable if:

there exists a serial execution order of the transactions

that yields the same results as the actual execution

(serial means one at a time -- no parallel execution)

(this definition should remind you of linearizability)

You can test whether an execution's result is serializable by

looking for an order that yields the same results.

for our example, the possible serial orders are

T1; T2

T2; T1

so the correct (serializable) results are:

T1; T2 : x=11 y=9 "11,9"

T2; T1 : x=11 y=9 "10,10"

the results for the two differ; either is OK

no other result is OK

the implementation might have executed T1 and T2 in parallel

but it must still yield results as if in a serial order

what if T1's operations run entirely between T2's two get()s?
 would the result be serializable?
 T2 would print 10,9
 but 10,9 is not one of the two serializable results!
 what if T2 runs entirely between T1's two adds()s?
 T2 would print 11,10
 but 11,10 is not one of the two serializable results!
 what if x's server does the increment but y's server fails?
 x=11 y=10 is not one of the serializable results!

Why serializability is popular

An easy model for programmers

They can write complex transactions while ignoring concurrency

It allows parallel execution of transactions on different records

a transaction can "abort" if something goes wrong

an abort un-does any record modifications

the transaction might voluntarily abort,

e.g. if the account doesn't exist, or y's balance is ≤ 0

the system may force an abort, e.g. to break a locking deadlock

some servers failures result in abort

the application might (or might not) try the transaction again

distributed transactions have two big components:

concurrency control (to provide isolation/serializability)

atomic commit (to provide atomicity despite failure)

first, concurrency control

correct execution of concurrent transactions

two classes of concurrency control for transactions:

pessimistic:

lock records before use

conflicts cause delays (waiting for locks)

optimistic:

use records without locking

commit checks if reads/writes were serializable

conflict causes abort+retry

called Optimistic Concurrency Control (OCC)

pessimistic is faster if conflicts are frequent

optimistic is faster if conflicts are rare

today: pessimistic concurrency control

next week: optimistic concurrency control (FaRM)

"Two-phase locking" is one way to implement serializability

2PL definition:

a transaction must acquire a record's lock before using it

a transaction must hold its locks until *after* commit or abort

2PL for our example

suppose T1 and T2 start at the same time

the transaction system automatically acquires locks as needed

so first of T1/T2 to use x will get the lock

the other waits until the first completely finishes

this prohibits the non-serializable interleavings

details:

each database record has a lock

if distributed, the lock is typically stored at the record's server

[diagram: clients, servers, records, locks]

(but two-phase locking isn't affected much by distribution)

an executing transaction acquires locks as needed, at the first use

add() and get() implicitly acquires record's lock

end_xaction() releases all locks

all locks are exclusive (for this discussion, no reader/writer locks)

the full name is "strong strict two-phase locking"
 related to thread locking (e.g. Go's Mutex), but easier:
 explicit begin/end_xaction
 DB locks automatically, on first use of each record
 DB unlocks automatically, at transaction end
 DB may automatically abort to cure deadlock

Why hold locks until after commit/abort?

why not release as soon as done with the record?

example of a resulting problem:

suppose T2 releases x's lock after get(x)

T1 could then execute between T2's get()s

T2 would print 10,9

oops: that is not a serializable execution: neither T1;T2 nor T2;T1

example of a resulting problem:

suppose T1 writes x, then releases x's lock

T2 reads x and prints

T1 then aborts

oops: T2 used a value that never really existed

we should have aborted T2, which would be a "cascading abort"; awkward

Two-phase locking can produce deadlock, e.g.

```
T1      T2
get(x)  get(y)
get(y)  get(x)
```

The system must detect (cycles? lock timeout?) and abort a transaction

Could 2PL ever forbid a correct (serializable) execution?

yes; example:

```
T1      T2
get(x)
        get(x)
        put(x,2)
put(x,1)
```

locking would forbid this interleaving

but the result (x=1) is serializable (same as T2;T1)

The Question: describe a situation where Two-Phase Locking yields higher performance than Simple Locking. Simple locking: lock *every* record before *any* use; release after abort/commit.

Next topic: distributed transactions versus failures

how can distributed transactions cope with failures?

suppose, for our example, x and y are on different "worker" servers

suppose x's server adds 1, but y's crashes before subtracting?

or x's server adds 1, but y's realizes the account doesn't exist?

or x and y both can do their part, but aren't sure if the other will?

We want "atomic commit":

A bunch of computers are cooperating on some task

Each computer has a different role

Want to ensure atomicity: all execute, or none execute

Challenges: failures, performance

We're going to develop a protocol called "two-phase commit"

Used by distributed databases for multi-server transactions

The setting

Data is sharded among multiple servers

Transactions run on "transaction coordinators" (TCs)

For each read/write, TC sends RPC to relevant shard server

Each is a "participant"

Each participant manages locks for its shard of the data

There may be many concurrent transactions, many TCs

TC assigns unique transaction ID (TID) to each transaction

Every message, every table entry tagged with TID
To avoid confusion

Two-phase commit without failures:

[time diagram: TC, A, B]

TC sends put(), get(), &c RPCs to A, B

The modifications are tentative, only to be installed if commit.

TC gets to the end of the transaction.

TC sends PREPARE messages to A and B.

If A is willing to commit,

A responds YES.

then A is in "prepared" state.

otherwise, A responds NO.

Same for B.

If both A and B say YES, TC sends COMMIT messages to A and B.

If either A or B says NO, TC sends ABORT messages.

A/B commit if they get a COMMIT message from the TC.

I.e. they write tentative records to the real DB.

And release the transaction's locks on their records.

A/B acknowledge COMMIT message.

Why is this correct so far?

Neither A or B can commit unless they both agreed.

What if B crashes and restarts?

If B sent YES before crash, B must remember (despite crash)!

Because A might have received a COMMIT and committed.

So B must be able to commit (or not) even after a reboot.

Thus participants must write persistent (on-disk) state:

B must remember on disk before saying YES, including modified data.

If B reboots, and disk says YES but no COMMIT,

B must ask TC, or wait for TC to re-send.

And meanwhile, B must continue to hold the transaction's locks.

If TC says COMMIT, B copies modified data to real data.

What if TC crashes and restarts?

If TC might have sent COMMIT before crash, TC must remember!

Since one worker may already have committed.

Thus TC must write COMMIT to disk before sending COMMIT msgs.

And repeat COMMIT if it crashes and reboots,

or if a participant asks (i.e. if A/B didn't get COMMIT msg).

Participants must filter out duplicate COMMITs (using TID).

What if TC never gets a YES/NO from B?

Perhaps B crashed and didn't recover; perhaps network is broken.

TC can time out, and abort (since has not sent any COMMIT msgs).

Good: allows servers to release locks.

What if B times out or crashes while waiting for PREPARE from TC?

B has not yet responded to PREPARE, so TC can't have decided commit

so B can unilaterally abort, and release locks

respond NO to future PREPARE

What if B replied YES to PREPARE, but doesn't receive COMMIT or ABORT?

Can B unilaterally decide to abort?

No! TC might have gotten YES from both,

and sent out COMMIT to A, but crashed before sending to B.

So then A would commit and B would abort: incorrect.

B can't unilaterally commit, either:

A might have voted NO.

So: if B voted YES, it must "block": wait for TC decision.

Note:

The commit/abort decision is made by a single entity -- the TC.

This makes two-phase commit relatively straightforward.
The penalty is that A/B, after voting YES, must wait for the TC.

When can TC completely forget about a committed transaction?
If it sees an acknowledgement from every participant for the COMMIT.
Then no participant will ever need to ask again.

When can participant completely forget about a committed transaction?
After it acknowledges the TC's COMMIT message.
If it gets another COMMIT, and has no record of the transaction,
it must have already committed and forgotten, and can acknowledge (again).

Two-phase commit perspective

Used in sharded DBs when a transaction uses data on multiple shards
But it has a bad reputation:
slow: multiple rounds of messages
slow: disk writes
locks are held over the prepare/commit exchanges; blocks other xactions
TC crash can cause indefinite blocking, with locks held
Thus usually used only in a single small domain
E.g. not between banks, not between airlines, not over wide area
Faster distributed transactions are an active research area.

Raft and two-phase commit solve different problems!

Use Raft to get high availability by replicating
i.e. to be able to operate when some servers are crashed
the servers all do the *same* thing
Use 2PC when each participant does something different
And *all* of them must do their part
2PC does not help availability
since all servers must be up to get anything done
Raft does not ensure that all servers do something
since only a majority have to be alive

What if you want high availability *and* atomic commit?

Here's one plan.
[diagram]
The TC and servers should each be replicated with Raft
Run two-phase commit among the replicated services
Then you can tolerate failures and still make progress
You'll build something like this to transfer shards in Lab 4
Next meeting's Spanner uses this arrangement